

# naalaithiran-project

October 15, 2023

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')
```

```
[ ]: dataset = pd.read_csv('/content/diabetes.csv')
```

```
[ ]: dataset.head()
```

```
[ ]: 
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[ ]: dataset.shape
```

```
[ ]: (768, 9)
```

```
[ ]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
```

```

1   Glucose          768 non-null   int64
2   BloodPressure    768 non-null   int64
3   SkinThickness    768 non-null   int64
4   Insulin          768 non-null   int64
5   BMI              768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age              768 non-null   int64
8   Outcome          768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB

```

```
[ ]: dataset.describe().T
```

```
[ ]:
```

	count	mean	std	min	25%	\
Pregnancies	768.0	3.845052	3.369578	0.000	1.00000	
Glucose	768.0	120.894531	31.972618	0.000	99.00000	
BloodPressure	768.0	69.105469	19.355807	0.000	62.00000	
SkinThickness	768.0	20.536458	15.952218	0.000	0.00000	
Insulin	768.0	79.799479	115.244002	0.000	0.00000	
BMI	768.0	31.992578	7.884160	0.000	27.30000	
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.078	0.24375	
Age	768.0	33.240885	11.760232	21.000	24.00000	
Outcome	768.0	0.348958	0.476951	0.000	0.00000	


	50%	75%	max
Pregnancies	3.0000	6.00000	17.00
Glucose	117.0000	140.25000	199.00
BloodPressure	72.0000	80.00000	122.00
SkinThickness	23.0000	32.00000	99.00
Insulin	30.5000	127.25000	846.00
BMI	32.0000	36.60000	67.10
DiabetesPedigreeFunction	0.3725	0.62625	2.42
Age	29.0000	41.00000	81.00
Outcome	0.0000	1.00000	1.00

```
[ ]: dataset.isnull().sum()
```

```
[ ]: Pregnancies      0
      Glucose          0
      BloodPressure    0
      SkinThickness    0
      Insulin          0
      BMI              0
      DiabetesPedigreeFunction 0
      Age              0
      Outcome          0
dtype: int64
```

```
[ ]: sns.countplot(x = 'Outcome',data = dataset)
```

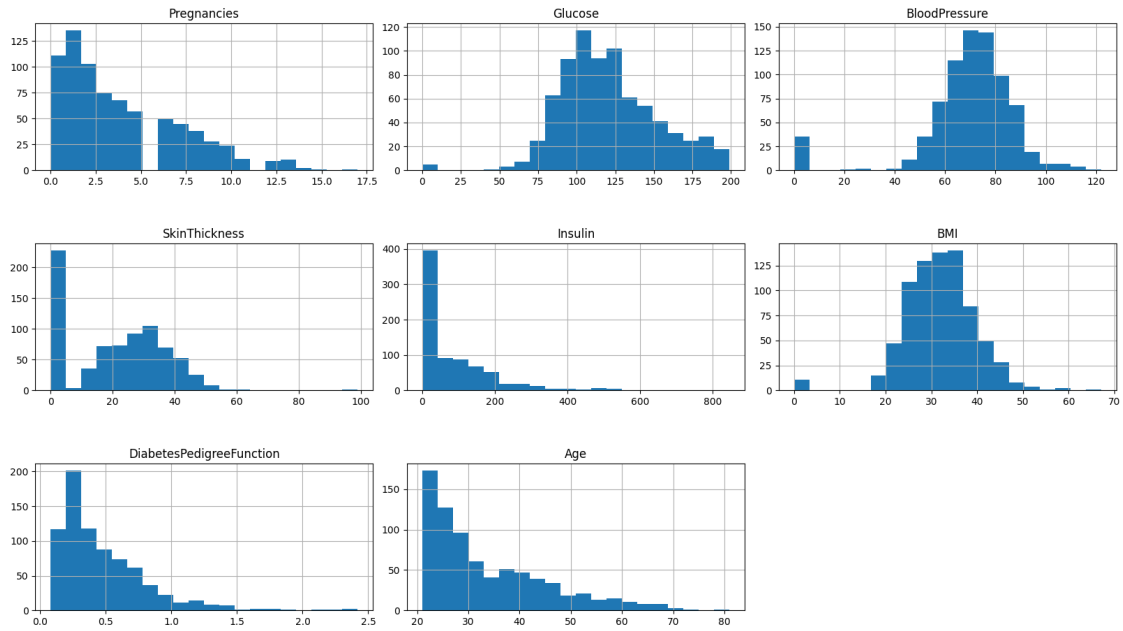
```
[ ]: <Axes: xlabel='Outcome', ylabel='count'>
```



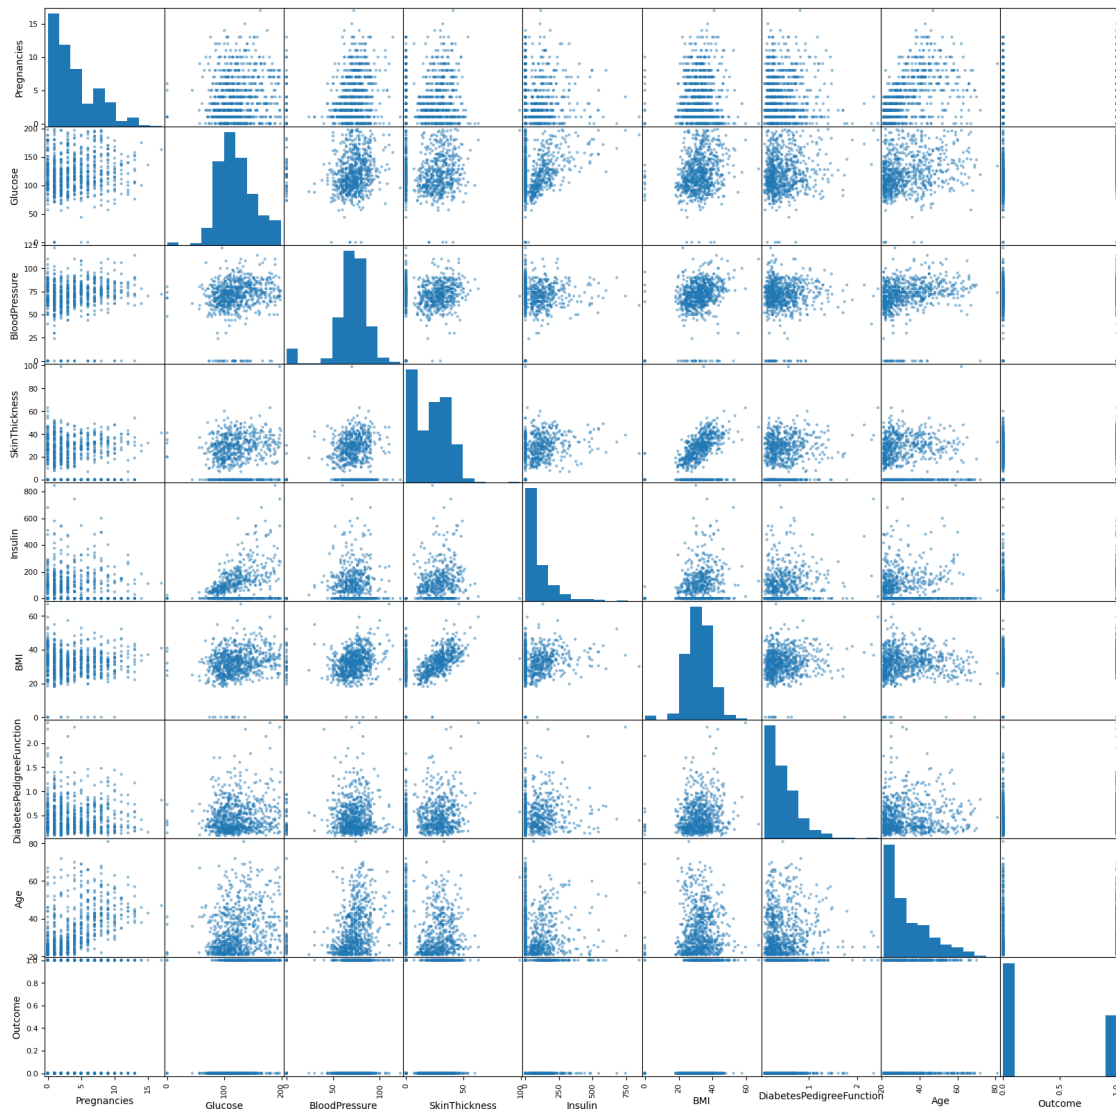
```
[ ]: import itertools

col = dataset.columns[:8]
plt.subplots(figsize = (20, 15))
length = len(col)

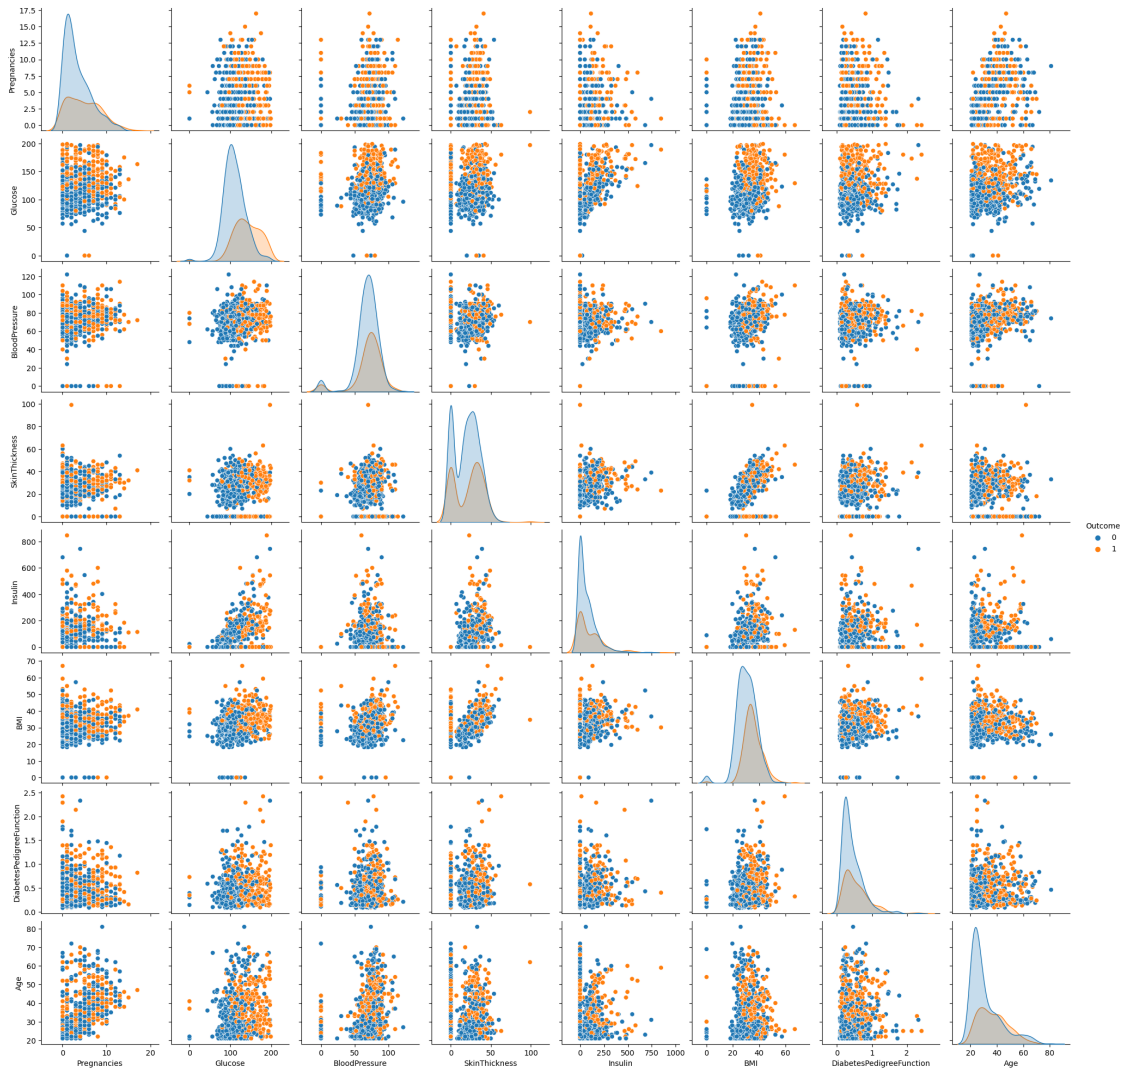
for i, j in itertools.zip_longest(col, range(length)):
    plt.subplot((length//2), 3, j + 1)
    plt.subplots_adjust(wspace = 0.1,hspace = 0.5)
    dataset[i].hist(bins = 20)
    plt.title(i)
plt.show()
```



```
[ ]: from pandas.plotting import scatter_matrix
scatter_matrix(dataset, figsize = (20, 20));
```



```
[ ]: sns.pairplot(data = dataset, hue = 'Outcome')
plt.show()
```



```
[ ]: pip install pandas
```

Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (1.5.3)


Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2023.3.post1)

Requirement already satisfied: numpy>=1.21.0 in /usr/local/lib/python3.10/dist-packages (from pandas) (1.23.5)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.1->pandas) (1.16.0)

```
[ ]: sns.heatmap(dataset.corr(), annot = True)
plt.show()
```



```
[ ]: dataset_new = dataset
```

```
[ ]: dataset_new[["Glucose", "BloodPressure", "SkinThickness", "Insulin", "BMI"]] =
↳ dataset_new[["Glucose", "BloodPressure", "SkinThickness", "Insulin", "BMI"]].
↳ replace(0, np.NaN)
```

```
[ ]: dataset_new.isnull().sum()
```

```
[ ]: Pregnancies      0
      Glucose         5
      BloodPressure   35
      SkinThickness   227
      Insulin        374
      BMI            11
```

```
DiabetesPedigreeFunction    0
Age                        0
Outcome                    0
dtype: int64
```

```
[ ]: dataset_new["Glucose"].fillna(dataset_new["Glucose"].mean(), inplace = True)
dataset_new["BloodPressure"].fillna(dataset_new["BloodPressure"].mean(),
    ↳inplace = True)
dataset_new["SkinThickness"].fillna(dataset_new["SkinThickness"].mean(),
    ↳inplace = True)
dataset_new["Insulin"].fillna(dataset_new["Insulin"].mean(), inplace = True)
dataset_new["BMI"].fillna(dataset_new["BMI"].mean(), inplace = True)
```

```
[ ]: dataset_new.describe().T
```

```
[ ]:
```

	count	mean	std	min	25%	\
Pregnancies	768.0	3.845052	3.369578	0.000	1.00000	
Glucose	768.0	121.686763	30.435949	44.000	99.75000	
BloodPressure	768.0	72.405184	12.096346	24.000	64.00000	
SkinThickness	768.0	29.153420	8.790942	7.000	25.00000	
Insulin	768.0	155.548223	85.021108	14.000	121.50000	
BMI	768.0	32.457464	6.875151	18.200	27.50000	
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.078	0.24375	
Age	768.0	33.240885	11.760232	21.000	24.00000	
Outcome	768.0	0.348958	0.476951	0.000	0.00000	

	50%	75%	max
Pregnancies	3.000000	6.000000	17.00
Glucose	117.000000	140.250000	199.00
BloodPressure	72.202592	80.000000	122.00
SkinThickness	29.153420	32.000000	99.00
Insulin	155.548223	155.548223	846.00
BMI	32.400000	36.600000	67.10
DiabetesPedigreeFunction	0.372500	0.626250	2.42
Age	29.000000	41.000000	81.00
Outcome	0.000000	1.000000	1.00

```
[ ]: from sklearn.preprocessing import MinMaxScaler
sc = MinMaxScaler(feature_range = (0, 1))
dataset_scaled = sc.fit_transform(dataset_new)
```

```
[ ]: dataset_scaled = pd.DataFrame(dataset_scaled)
```

```
[ ]: X = dataset_scaled.iloc[:, [1, 4, 5, 7]].values
Y = dataset_scaled.iloc[:, 8].values
```



```
[ ]: from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.20,
↳ random_state = 42, stratify = dataset_new['Outcome'] )
```

```
[ ]: print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("Y_train shape:", Y_train.shape)
print("Y_test shape:", Y_test.shape)
```

```
X_train shape: (614, 4)
X_test shape: (154, 4)
Y_train shape: (614,)
Y_test shape: (154,)
```


```
[ ]: from sklearn.linear_model import LogisticRegression
logreg = LogisticRegression(random_state = 42)
logreg.fit(X_train, Y_train)
```

```
[ ]: LogisticRegression(random_state=42)
```

```
[ ]: from sklearn import metrics
from sklearn.neighbors import KNeighborsClassifier

X_axis = list(range(1, 31))
acc = pd.Series()
x = range(1,31)

for i in list(range(1, 31)):
    knn_model = KNeighborsClassifier(n_neighbors = i)
    knn_model.fit(X_train, Y_train)
    prediction = knn_model.predict(X_test)
    acc = acc.append(pd.Series(metrics.accuracy_score(prediction, Y_test)))
plt.plot(X_axis, acc)
plt.xticks(x)
plt.title("Finding best value for n_estimators")
plt.xlabel("n_estimators")
plt.ylabel("Accuracy")
plt.grid()
plt.show()
print('Highest value: ',acc.values.max())
```



Highest value: 0.7857142857142857

```
[ ]: from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors = 24, metric = 'minkowski', p = 2)
knn.fit(X_train, Y_train)
```

```
[ ]: KNeighborsClassifier(n_neighbors=24)
```

```
[ ]: from sklearn.svm import SVC
svc = SVC(kernel = 'linear', random_state = 42)
svc.fit(X_train, Y_train)
```

```
[ ]: SVC(kernel='linear', random_state=42)
```

```
[ ]: from sklearn.naive_bayes import GaussianNB
nb = GaussianNB()
nb.fit(X_train, Y_train)
```

```
[ ]: GaussianNB()
```

```
[ ]: from sklearn.tree import DecisionTreeClassifier
dectree = DecisionTreeClassifier(criterion = 'entropy', random_state = 42)
dectree.fit(X_train, Y_train)
```

```
[ ]: DecisionTreeClassifier(criterion='entropy', random_state=42)
```

```
[ ]: from sklearn.ensemble import RandomForestClassifier
ranfor = RandomForestClassifier(n_estimators = 11, criterion = 'entropy',
    random_state = 42)
ranfor.fit(X_train, Y_train)
```

```
[ ]: RandomForestClassifier(criterion='entropy', n_estimators=11, random_state=42)
```

```
[ ]: Y_pred_logreg = logreg.predict(X_test)
Y_pred_knn = knn.predict(X_test)
Y_pred_svc = svc.predict(X_test)
Y_pred_nb = nb.predict(X_test)
Y_pred_dectree = dectree.predict(X_test)
Y_pred_ranfor = ranfor.predict(X_test)
```

```
[ ]: from sklearn.metrics import accuracy_score
accuracy_logreg = accuracy_score(Y_test, Y_pred_logreg)
accuracy_knn = accuracy_score(Y_test, Y_pred_knn)
accuracy_svc = accuracy_score(Y_test, Y_pred_svc)
accuracy_nb = accuracy_score(Y_test, Y_pred_nb)
accuracy_dectree = accuracy_score(Y_test, Y_pred_dectree)
accuracy_ranfor = accuracy_score(Y_test, Y_pred_ranfor)
```

```
[ ]: print("Logistic Regression: " + str(accuracy_logreg * 100))
print("K Nearest neighbors: " + str(accuracy_knn * 100))
print("Support Vector Classifier: " + str(accuracy_svc * 100))
print("Naive Bayes: " + str(accuracy_nb * 100))
print("Decision tree: " + str(accuracy_dectree * 100))
print("Random Forest: " + str(accuracy_ranfor * 100))
```


Logistic Regression: 72.07792207792207  
K Nearest neighbors: 78.57142857142857  
Support Vector Classifier: 73.37662337662337  
Naive Bayes: 71.42857142857143  
Decision tree: 68.18181818181817  
Random Forest: 75.97402597402598

```
[ ]: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(Y_test, Y_pred_knn)
cm
```

```
[ ]: array([[87, 13],
          [20, 34]])
```

```
[ ]: sns.heatmap(pd.DataFrame(cm), annot=True)
```

```
[ ]: <Axes: >
```



```
[ ]: from sklearn.metrics import classification_report
print(classification_report(Y_test, Y_pred_knn))
```

	precision	recall	f1-score	support
0.0	0.81	0.87	0.84	100
1.0	0.72	0.63	0.67	54
accuracy			0.79	154
macro avg	0.77	0.75	0.76	154
weighted avg	0.78	0.79	0.78	154